

Pitch Q&A Preparation

Team DataBots | Problem Statement 26154 (NTRO)

Security & Privacy (The NTRO Priority)

Q: How is the data handled? Cloud vs. On-Prem?

A: We designed this with strict data sovereignty in mind. The architecture supports on-premise deployment using open-weight models (like Llama 3 or Mistral) so sensitive NTRO data never leaves the internal network. If deployed on the cloud, it uses secure enterprise endpoints where data is strictly isolated and zero data is used for model training.

Q: What is your data retention policy?

A: Our platform operates with a 'process and purge' architecture. Documents are temporarily held in memory during the generation phase and are immediately purged from the system once the output is exported, leaving no residual footprint.

AI Accuracy & Hallucinations

Q: How does your platform prevent AI hallucinations?

A: We tackle this in two ways. First, we use strict system prompts that constrain the AI to only use the information in the provided document (similar to a RAG approach). Second, we set the model's 'temperature' very low, prioritizing factual accuracy and precision over creative guessing.

Q: Is there a review process before outputs are finalized?

A: This is a co-pilot, not an autopilot. Every piece of generated content is rendered in an editable dashboard. The human operator remains the final gatekeeper, reviewing and approving the text before it is copied or exported.

Technical Architecture

Q: How does the system handle multiple output formats requested at once?

A: To ensure speed, the system processes multiple outputs concurrently. We use asynchronous API calls, meaning that if you request a tweet, a PDF summary, and a video script, the engine generates them simultaneously rather than making you wait for them one by one.

Q: How do you extract text from images and scanned PDFs?

A: We use a multi-step extraction pipeline. For text documents and standard PDFs, we parse the text directly. For images or scanned documents, we route the file through an OCR (Optical Character Recognition) tool or a Vision model to extract the text and layout context before sending it to the generation engine.

Product & User Experience

Q: How does the system customize tone and style for different formats?

A: We use dynamic prompt injection. Under the hood, each output format has a pre-engineered system prompt (e.g., 'Act as a senior cybersecurity analyst writing an advisory'). Furthermore, admins can add their own custom templates and define the exact tone rules required for their specific agency.